

A Model of Criminal Justice Discretion*

Tom S. Clark
Professor of Political Science
Stanford University
Stanford, CA, USA
tsclark@stanford.edu
ORCID: 0000-0002-3552-9161

August 24, 2026

Abstract

Criminal prosecutors wield tremendous power in the American legal system. However, their ability to direct prosecution against criminal suspects relies on coordination with the police who investigate crimes and make arrests. That relationship has become more prominent with the rise of reform prosecutors who propose to shift traditional law-enforcement priorities. I develop a model of police-prosecutor interaction built around an informational asymmetry that past research on prosecution has not analyzed. Police hold private information about the culpability of an arrested suspect, but they also have their own preferences over how that suspect is charged. I ask how a prosecutor's interest in leniency affects whether police can credibly communicate what they know. Two findings stand out. First, alignment between police and prosecutors is observationally equivalent across very different underlying relationships. Apparent cooperation can reflect either credible communication or its absence, so the alignment we observe cannot by itself reveal whether police information is reaching prosecutors. Second, the court's willingness to acquit an over-charged suspect, ordinarily treated as waste, is what makes truthful communication possible, because the prospect of acquittal disciplines the allegations police bring. These results caution against reading cooperation between police and prosecutors as evidence of well-functioning justice, and they bear on how we evaluate criminal justice reform.

Keywords: Prosecutorial discretion · Police–prosecutor relations · Cheap talk · Asymmetric information · Criminal justice reform

Funding: This research was not supported by funding.

*I thank Sam England, Sean Gailmard, Ryan Hübert, John Patty and Ian Turner for helpful comments and discussions.

1 Introduction

In recent years, Americans have debated whether and how to reform the criminal justice system. The stakes are high: criminal justice is how government both protects individual rights and infringes on them, and the system is staffed by officials whose discretion is largely immune from meaningful check. Recognizing that significance, political scientists have turned anew to policing, prosecution, and the carceral state. In many ways, prosecutors are the linchpin in this system. They sit between policing and judicial disposition, wield resources that dwarf those of defendants, and enjoy remarkable discretion over how to charge. Often, though not always, they are elected, and they face complex incentives in handling the cases police bring them.¹

The recent debate has centered on so-called “reform” prosecutors. When a prosecutor adopts priorities at odds with the police, trust between the two erodes: in cities from San Francisco to Chicago to Burlington, Vermont, police and their unions have cast reform prosecutors as “soft on crime” and, in some cases, reduced cooperation or threatened to “stand down” (e.g., Thompson and Murphy 2024). That breakdown, though, has a particularly asymmetric nature. Police protest prosecutors who decline to charge as aggressively as police want, but we do not observe police protesting prosecutors for being *too tough*. This directional asymmetry is at the heart of the strategic problem I model.

How does trust between police and prosecutors shape the implementation of criminal law? The question bears directly on American political institutions, since these are the officials who apply most law at the individual level and form the contact point between enforcement and adjudication. I develop a simple model of their collaboration. The model turns on a specific informational problem. Through their street-level work, police accumulate private, character-based knowledge of a suspect’s culpability, knowledge that is tacit and difficult to convert into a form the prosecutor can verify. The question I ask is when police can credibly

¹The model applies most directly to the elected prosecutors of most state and local systems; appointed and federal prosecutors face different pressures, though career concerns may generate similar dynamics.

convey what they know, and what coordination between police and prosecutors looks like when they cannot.

I present three findings with implications for what we can learn about police-prosecutor relations from what we observe. First, the results show that very different relationships between police and prosecutors are observationally equivalent. What looks like cooperation can arise when the two credibly communicate to coordinate on appropriate charges, and it can arise when they do not trust each other and the quality of justice suffers. The alignment we usually observe, police allegations and prosecutorial charges moving together, is consistent with both. Thus, whether a reform is working depends on which relationship underlies the cooperation we see, rather than the mere existence of cooperation. Learning which regime we are in requires looking at other features of the system, especially the cases where police and prosecutors diverge.

The second finding highlights a counter-intuitive role for courts. In the usual economic analysis of criminal justice, acquittal is costly: a case the system failed to win, or deterrence left uncollected. However, in this model, acquittal is productive. Convicting a defendant of a more serious charge than his conduct warrants is harder than convicting him of a lesser one, and that difficulty can induce police to communicate honestly. The risk of a case leading to acquittal because a defendant is over-charged creates an incentive for a police officer to avoid exaggerating a defendant's culpability and, with that, constraint on the kinds of allegations he is willing to make. The prospect of acquittal, ordinarily treated as waste, is what lets police information travel credibly to the prosecutor. Relatedly, the value police and prosecutors place on maintaining trust can degrade their communication rather than improve it. As the cost of public disagreement grows, prosecutors become more reluctant to deviate from police allegations, and police exploit that reluctance.

The third finding is that changes to one part of the criminal justice system propagate through the others in specific ways the model makes precise. Debates about reforming criminal justice to avoid over-punishment cannot focus on prosecutors alone. They have

to account for the interaction among prosecutors, police, and courts, and that interaction can complicate the goals of prosecutorial reform in ways that are easy to miss when each institution is studied on its own.

2 The Law Enforcement Community

2.1 Prosecutors, Police, and their Incentives

While prosecutors play a significant role in the politics of crime and punishment, they have remained largely understudied in political science (Gordon and Huber 2009). Traditionally, scholars have approached prosecutors as relatively simple actors who seek to maximize sentences or convictions (e.g., Bar-Gill and Gazal Ayal 2006, Landes 1971); they are seen as agents of social welfare who use discretion to maximize compliance with the law and ensure efficient deterrence (e.g., Baker and Mezzetti 2001, Easterbrook 1983). That approach overlooks two political features of the criminal justice system. First, prosecutors are elected officials, and they face strong incentives to secure convictions not just for social welfare reasons but because of the electoral benefits that come from convictions (e.g., Gordon and Huber 2002, 2009). Second, prosecutors are part of an institutional matrix that includes police, judges, prison systems, and law-makers, and they face potentially cross-cutting pressures when deciding how and when to enforce the law (e.g., Clark 2026, Davis 2007, Romero 2019, Starr and Rehavi 2013).

I start from those observations to model prosecutorial discretion as part of an institutional setting with many of the classic hallmarks of bureaucratic decision-making. However, in contrast to standard theories of bureaucratic decision-making that focus on hierarchy or informational asymmetries between principals and agents (e.g., Makris 2006, Miller 1977, Steunenberg 1996, Weingast 1984), I propose a novel informational problem. Whereas classic models pair a political principal with a subordinate bureau, in this paper I model two *peer* bureaucrats, police and prosecutor, neither one the other's superior. In the model, they

must coordinate charging decisions under asymmetric information, with no principal able to verify the officer’s private assessment. The communication friction I model is thus a specific instance of the information and discretion problems at the center of public-choice theories of bureaucracy, relocated to a horizontal relationship between two enforcement agencies.

Research that situates prosecution in its political environment has concentrated on plea bargaining, where the friction is informational asymmetry between prosecutor and defendant, and on how career and electoral incentives shape discretion (e.g., Gordon 2009, Gordon and Huber 2002, Little and Simpson 2024). Shifts in prosecutorial incentives measurably affect charging, incarceration, and pleas (e.g., Agan, Doleac and Harvey 2023, Hessick 2021, Ouss 2020), and many scholars conclude that “prosecutors rule the criminal justice system” (Clark Forthcoming, Rakoff 2016).

What that work largely sets aside is the relationship between prosecutors and the police, who must cooperate because prosecutors rely on officers to identify suspects and build cases (for discussion, see Garoupa 2012, Romero 2019). It is the public sense that police have responded to reform prosecutors by “standing down” that motivates much of the recent critique of those prosecutors. Why, and under what conditions, might reform prosecutors undermine the effectiveness of policing itself?

There are several forms of “standing down.” It can be a retreat from discretionary, self-initiated enforcement (e.g, fewer stops, searches, and arrests) of the kind measured in the “de-policing” that followed high-profile accountability episodes (e.g., Shjarback et al. 2017); slower or less thorough investigative work; de-prioritizing the categories of case a reform prosecutor is known to decline; and public messaging, often through police unions, that casts the prosecutor as soft on crime. In the confrontations between reform prosecutors and police in cities such as Chicago and Philadelphia, these responses have appeared together and have materially undercut prosecutors’ agendas (e.g., Kohler 2023, Romero 2019). The common element that is central to the argument here is officers withholding discretionary effort and cooperation rather than openly defying the prosecutor.

These conflicts highlight an often overlooked aspect of the relationship between prosecutors and police. Officers exercise wide discretion, and while individual biases shape their decisions (e.g., Ba et al. 2021, Clark et al. 2023, Clark, Owens and Glynn 2025, Hübert and Little 2023, Knox, Lowe and Mummolo 2020), they also respond to political incentives and their environment (e.g., Clark 2024, Mummolo 2018) and are motivated by a desire to see justice done. In what Wilson calls the “Watchman” form of policing, officers judge “the seriousness of infractions less by what the law says about them than by their immediate and personal consequences. . .” (Wilson 1968, 141). That is, police develop a sense of an individual criminal suspect’s culpability—the degree to which that person is truly responsible for, and likely to persist in, criminal behavior—in a way that is detached, to some extent, from the specific legal implications of his or her actions. And, importantly, police value matching their perception of a suspect’s underlying culpability with the way that person is handled by the criminal justice system. Indeed, this character-based dimension of culpability is not unique to police: judges and juries respond to it as well, weighing demeanor, community standing, and the circumstances of an offense alongside its formal elements, even when the formal evidence is held constant (e.g., Pennington and Hastie 1992).

Doing that work requires close, intimate knowledge of the people over whom police have jurisdiction. Criminal justice scholars point to co-production—the process by which police work with communities to provide public safety—as a mechanism by which police develop informational advantages that allow them to better assess the culpability of individuals they encounter, though that process also endows police with the power to manipulate the criminal justice system and advance their own biases (e.g., Brewer and Grabosky 2014). And because police are just one part of that system, they also value cooperation and trust with the prosecutor. a coordination so close it has been called an “unholy union” (e.g., Romero 2019). Police value their ability to influence prosecutors’ decisions by bringing cases to them and framing those cases in a particular way. In the model I develop below, police value cooperation with the prosecutor, particularly cooperation that aligns charging decisions with

the suspect’s underlying culpability.

2.2 Street-Level Bureaucrats and City-Level Decision-Makers

The interaction between police and prosecutors entails, then, two actors with wide discretion. They individually have significant power, but the criminal justice system functions as a product of their joint efforts. While they share some common goals, their objectives can diverge. The recent conflicts between police and reform prosecutors (and episodes of standing down) highlight the challenge that arises when they do. A pressing question has emerged: can police and prosecutors cooperate when a prosecutor does not share the same law enforcement priorities as the officers themselves?

The model I build centers on an asymmetric version of a vulnerable trust relationship between police and prosecutors. The episodes of conflict generally run in one direction: police criticize, protest, and sometimes operationally withdraw from prosecutors who decline to charge at the level they want (Romero 2019, Thompson and Murphy 2024), while the reverse—protest against prosecutors who charge too aggressively—is not something we tend to observe. That is, trust between police and prosecutors tends to be sensitive to prosecutors “going easy” on criminal defendants but not necessarily prosecutors “throwing the book” at them. The model I build directly captures this asymmetry.

The key tension I analyze is a dynamic that practitioners of criminal justice understand but research on law enforcement has yet to fully integrate. By virtue of their day-to-day, street-level interactions with individuals, police officers accumulate relevant information about the true culpability of an individual arrested for a crime. There are two elements that motivate this dynamic. The first relates to the actual substantive, qualitative information the police possess by way of their daily interactions with the public. The second relates to the structural, inherent difficulty of communicating that detailed knowledge to a prosecutor who is handling a case.

What police know that prosecutors do not. Police accumulate at least four kinds of information that prosecutors typically lack. The first is neighborhood-grounded recidivism information—knowledge of whether an arrest is the latest in a pattern of similar incidents at a given address or involving the same individuals, much of which never produces an arrest and so never appears in the formal criminal-history record that travels with the case to the prosecutor’s office (e.g., Kohler-Hausmann 2018). The second is informal social-network information: who in a neighborhood is feared, whom victims and witnesses regard as predatory but will not testify against on the record, who maintains standing relationships with other actors in a local illicit economy (e.g., Stuart 2016). The third is interactional information from the arrest encounter itself. Examples include demeanor, statements made off-record, body language, what the suspect did or did not volunteer in moments not captured by formal reports or body-worn cameras. The fourth, and analytically most interesting, is information generated from cumulative interactions over time, which officers integrate into a gestalt assessment but cannot easily decompose into discrete evidentiary claims, which is the central object of the literature on tacit knowledge in street-level bureaucracy (e.g., Lipsky 1980, Maynard-Moody and Musheno 2003). This gestalt quality is an instance of what Polanyi (1966) termed *tacit knowledge* and of Hayek’s (1948) broader point that much decision-relevant knowledge is dispersed and local, tied to the particular circumstances in which it is acquired. What officers learn from their day-to-day work that allows them to make an assessment of a criminal suspect’s culpability is exactly this kind of information: he can make use of it on the ground but cannot readily transfer it to a prosecutor who was not there to make her own assessment in the first place.

Why this information is costly to credibly communicate. This last part brings us to the key feature that leads to the strategic tension I model. Namely, the police have substantive knowledge about the criminal suspects that they cannot easily communicate. Three mechanisms documented in the institutional literature support that assumption. First, the

information is gestalt-like, synthesized impression rather than a list of discrete facts. Converting it into evidence-grade assertions imposes articulability costs on the officer that are not present for formal evidentiary information (e.g., Maynard-Moody and Musheno 2003). Second, putting reputation-based assessments into writing exposes the officer to professional and legal risk: such documents are discoverable in subsequent proceedings, can be the basis for cross-examination, and can be cited as evidence of bias (e.g., Richman 2003). Third, the institutional architecture of case handoff involves written reports rather than sustained conversations, time-pressured intake at the prosecutor’s office, formal evidentiary protocols. That nature structurally limits the bandwidth available for transmitting nuanced contextual assessments, even when both parties would prefer richer communication (e.g., Kohler-Hausmann 2018, Richman 2003). Together, these mechanisms explain why an officer’s private information may persist as private even when both the officer and the prosecutor would prefer a richer exchange.

Underlying all three mechanisms is a single feature: the officer’s assessment is not verifiable. If the prosecutor (or defense lawyer) could check it cheaply, none of the three would bind—a verifiable claim needs neither laborious articulation nor reputational cover, and could pass through even a narrow channel. It is precisely because such assessments cannot be verified that most well-functioning legal systems limit their use as a basis for conviction. Admitting unverifiable character judgments as evidence would hand officers discretionary conviction power with no guarantee they would exercise it truthfully. That is the danger I model below. Focusing on an officer with private, uncheckable information and his own stake in how the suspect is charged allows me to isolate the credibility of communication rather than the mere divergence of interests.

I model this institutional interaction with a simple binary, cheap-talk setting. And, as I show in the analysis, cheap talk fails to fully reveal the officer’s private information not because the underlying information is unobservable in principle but because the institutional message space is constrained in the ways just described. Prosecutors and police therefore

rely on trust in order to coordinate on how to handle individuals arrested, even individuals arrested under similar circumstances for similar crimes. This observation is critical for the results I present.

In particular, the model focuses on how the informational asymmetry between police and prosecutors interacts with their competing incentives. A key question is how a prosecutor’s use of discretion about charging criminal suspects affects the ability to work cooperatively with police. The model sets aside relevant actors, most notably civilians themselves as well as judges. In typical local criminal courts, judges defer to a large extent to prosecutors, seemingly due to prosecutors’ greater resources and deep familiarity with the case. That deference is itself an equilibrium phenomenon, and future models may extend the logic here to incorporate judicial decision-making directly, including the question, which the literature has yet to address, of how judges incorporate into their deference any reliance prosecutors place on information from the police themselves.

3 The Model

Players, actions, and sequence of play. There are two players: an Officer and a Prosecutor, denoted O and P , respectively. There is also a non-strategic player—the suspect—who makes no moves and has no decisions or payoffs.² The model revolves around how the Officer and the Prosecutor behave after the Officer has arrested the suspect. First, Nature selects a type for the suspect, $\theta \in \{L, H\}$, and an Officer type, $\beta \in \mathbf{R}$. The type θ represents the suspect’s *culpability* in the character-based sense developed in Section 2.1: the dimension of his case that would lead a judge or juror to want to treat him more or less harshly, holding constant the formal legal elements of the offense charged. (I will sometimes refer to $\theta = H$ as a “high type” suspect, meaning highly culpable in this character-based sense, and $\theta = L$ as a “low type” suspect. The terminology should not be confused with the formal severity of the charged offense, which is a separate dimension.) Let β denote the officer’s taste for

²Investigating incentives for that player is a valuable opportunity for future modeling.

severity, his preference for (or against) punitive charges. The parameter occupies the same structural position as the prosecutor’s electoral benefit η introduced below: it is a payoff the officer collects on the prosecutor’s charging decision, not a distortion of what he perceives or reports about the suspect. For that reason I describe β as a taste rather than a bias.

The Officer observes θ , but the Prosecutor does not. I assume that β is common knowledge. The Officer then decides what level of allegations to bring against the suspect, $a(\theta) \in \{L, H\}$. The Prosecutor receives the allegations from the Officer and decides how to charge the suspect in court, $c \in \{L, H\}$.

A disposition, $d \in \{0, 1\}$, is then probabilistically determined in court, which I assume is a function of the charges, c , and the suspect’s type, θ . Specifically, I assume that the suspect will be convicted with certainty except in one situation. If the Prosecutor chooses $c = H$ but the suspect is a low type—i.e., $\theta = L$ —then the suspect is found not guilty ($d = 0$) with probability $\pi \in (0, 1)$. The court does not observe θ directly; rather, the dependence of the disposition on θ reflects the fact that the same character-level factors captured by θ (e.g., prior pattern, demeanor, the circumstances surrounding the offense) surface during adjudication, through testimony, cross-examination, and the texture of the case itself, and shape how triers of fact respond to the charges.

The parameter π captures the court’s relative preference between two kinds of error: over-convicting low-type defendants on serious charges (which it permits with probability $1 - \pi$) and under-convicting them entirely (which it accepts with probability π). A high- π court is one that treats the character-level information that emerges at trial as grounds for reasonable doubt about the most serious charges, A court averse to attaching a high-severity label to a low-culpability defendant. More broadly, one might interpret π as capturing the disposition tendencies of the adjudication process as a whole, including the effectiveness of defense counsel in surfacing mitigating facts. While stark, this modeling choice focuses analytic attention on the disposition margin that is the source of conflict between the prosecutor and

the police: the risk that overcharging a defendant ends in acquittal.³

Preferences. The players’ utilities are summarized in Table 1. A number of the assumptions I make about players’ preferences diverge from a standard set-up, and so I discuss those in more substantive terms. I assume the Officer has preferences over four elements. First, the Officer has a taste for justice. Specifically, he prefers that the Prosecutor file charges that match the suspect’s true type and receives $\gamma > 0$ whenever his allegations match the defendant’s type and those allegations are adopted by the Prosecutor. That is, he receives positive utility from aligning Officer and Prosecutor choices with the defendant’s type. Second, I assume that the Officer receives utility associated with his type, $\beta \in \mathbf{R}$, whenever $c = H$. Because β can be positive or negative, the Officer can have a preference either for more serious charges or more lenient charges for the suspect. Notably, this source of utility depends on the Prosecutor’s choice, independent of the ultimate outcome for the defendant. I make this assumption to capture the political dynamic motivating the model—police have preferences about how prosecutors treat defendants. Third, I assume the Officer prefers convictions, independent of the charges. Specifically, the Officer pays a cost $\phi > 0$ whenever the suspect is not convicted. Finally, I assume the Officer pays a cost $\delta > 0$ whenever the Prosecutor reduces the Officer’s allegations, which reflects a loss in trust between the Officer and the Prosecutor.

I also assume the Prosecutor’s preferences capture four elements. First, the Prosecutor derives an electoral benefit from successful prosecutions. Let $\eta > 0$ denote the differential between securing a conviction on a “high type” charge (i.e., $c = H$) as opposed to a “low type” charge. Specifically, she receives η from a conviction of charges $c = H$. Convictions of more serious charges yield a stronger electoral benefit than do convictions of more lenient charges. Second, I assume the Prosecutor receives $\lambda \geq 0$ if she chooses $c = L$ when $\theta = L$.

³Charge severity here is the conviction label, not the sentence: the overcharge case ($c = H$, $\theta = L$) turns on whether the high-severity label attaches. Sentencing discretion, where some mitigating information also operates, is collapsed into this margin because the documented conflicts are about charging, not sentencing. The model also has no innocent type, so it sets aside conviction of the innocent, a limitation I discuss below.

	$c = H$	$c = L$
$a = H$	$\gamma + \beta, \eta$	$-\delta, -\delta$
$a = L$	$\beta, \eta - \epsilon$	$0, 0$

(a) $\theta = H$

	H	L
H	$\beta - \pi\phi, (1 - \pi)\eta$	$-\delta, \lambda - \delta$
L	$\beta - \pi\phi, (1 - \pi)\eta - \epsilon$	γ, λ

(b) $\theta = L$

Table 1: Players' payoffs for each type of suspect, θ .

Notes: Each cell reports (Officer payoff, Prosecutor payoff). Within each subtable, rows index the Officer's allegation $a \in \{H, L\}$ and columns index the Prosecutor's charge $c \in \{H, L\}$. The parameters are: $\gamma > 0$, the Officer's taste for matching charges to the suspect's type; $\beta \in \mathbf{R}$, the Officer's taste for severity; $\phi > 0$, the Officer's cost of an acquittal; $\delta > 0$, the trust cost paid when the Prosecutor downgrades the Officer's allegation; $\eta > 0$, the Prosecutor's electoral benefit from a high-type conviction; $\lambda \geq 0$, the Prosecutor's benefit from correctly charging a low-type suspect with a low-type charge (her interest in avoiding over-charging); $\epsilon > 0$, the Prosecutor's effort cost of escalating charges beyond a low-type allegation; and $\pi \in (0, 1)$, the probability the court acquits a low-type suspect charged with a high-type crime. Each parameter is defined in the text below.

This parameter measures, then, the prosecutor's interest in avoiding over-charging a suspect. (One could also include a parallel parameter to measure the prosecutor's benefit from correctly charging a high-type suspect with high-type charges. As I show in the appendix, that parameter does not affect the analysis of the model.) As a result of these assumptions, the penalty to the Prosecutor for acquittal—not convicting the defendant—is a function of η and λ . The prosecutor does not receive η in the event of an acquittal, but that can be offset by a large enough λ . Third, the Prosecutor has a distaste for effort. If the Officer presents a less serious allegation ($a = L$), the Prosecutor pays a cost $\epsilon > 0$ if she wants to charge a more serious crime, $c = H$, representing the work she must do to enhance the charges and build a bigger case. Finally, like the Officer, the Prosecutor has a taste for trust between her office and the Officer and pays a cost $\delta > 0$ whenever she reduces the Officer's allegations—i.e., $a = H$ but $c = L$.⁴

As will become evident in the analysis below, Table 1 shows that the primary tension between the Officer and the Prosecutor occurs when the suspect is a low-type. It is under this condition that the Officer is most likely to want high-type charges under conditions where the Prosecutor wants low-type charges. That potential conflict of interests is at the heart of the model.

⁴The model could index the δ parameter so that the Prosecutor and Officer have distinct costs for losing trust, and all of the results here would hold.

Although the model carries several parameters, the analysis is organized around a single substantive dimension of preference divergence—a police officer who leans toward more serious charges facing a prosecutor with a meaningful interest in leniency—and a small number of composite quantities do all of the work. The focal case throughout is $\beta > \gamma > 0$: an officer whose taste for severity outweighs his taste for matching charges to types. (When $\beta < \gamma$, including any officer who leans against serious charges, the officer has no incentive to misrepresent and the communication problem dissolves; the analysis below notes these cases where relevant.) On the Officer’s side, behavior turns on the ratio $(\beta - \gamma) / \phi$ —the taste for severity net of the taste for justice, scaled by the aversion to acquittal—which determines how willing the officer is to gamble on serious charges against a low type. On the Prosecutor’s side, behavior turns on the leniency interest λ relative to the conviction payoff η and the effort cost ϵ . The trust parameter δ enters both players’ calculations and shifts the thresholds at which these comparisons bite.

Information. There is one element of uncertainty in the model—the Prosecutor’s uncertainty about the suspect’s type, θ . I assume that the Prosecutor has a prior belief, $\Pr(\theta = H) = p \in (0, 1)$, and updates her belief according to Bayes’ Rule where appropriate.

Notable modeling choices. First, I treat the trial outcome as an exogenous, commonly known probability rather than an adversarial contest; π is the court’s equilibrium disposition tendency—the rate at which adjudication declines to attach serious charges to low-culpability defendants—which isolates the officer–prosecutor communication problem. The case resolves at trial rather than by plea because that problem is logically prior to any bargaining the charges frame.

Second, preferences attach to the Prosecutor’s observable *charge* rather than the court’s eventual disposition: β and δ load on the charge, while ϕ and η capture the players’ stakes in the disposition.

Third, the type θ is the officer’s soft, contextual assessment of culpability (Section 2.2),

not the formal record the prosecutor also observes; its subjectivity is what makes credible communication hard. Because θ also shifts the odds of conviction—high types convicted with certainty, over-charged low types sometimes acquitted—the officer’s assessment is correlated with the strength of the case, which is why the prosecutor attends to it.⁵ Fourth, the Prosecutor’s action is binary—escalate or not. Costless alternative charging would dominate and drain the decision of information and political content, so the binary action is the reduced form of a world in which hedging is costly enough to preserve the charge as a commitment.⁶

Fifth, the type is exogenous and the game begins after arrest. The arrest pool is itself endogenous to prosecutorial behavior. Standing down’ episodes are arrest-rate responses, endogenizing it is left to future work.

Sixth, the trust cost δ is paid only when the Prosecutor *downgrades* an allegation, never when she escalates, formalizing the directional pattern of Section 2.2: a downgrade publicly repudiates the officer’s committed characterization in a way escalation does not.

Seventh, I focus on an Officer with a positive and sizable taste for severity, $\beta > \gamma > 0$: substantively Wilson’s (1968) “Watchman” officer, who values punitive charges beyond their fit with the record (see also Brewer and Grabosky 2014), and analytically the only interesting region, since for $-\gamma < \beta < \gamma$ the Officer communicates truthfully for any court (Corollaries 3–4). That police stand down only against prosecutors seen as too *lenient*, never too *tough* (e.g., Romero 2019), is itself evidence for the restriction.

4 Analysis

I analyze the model in four steps. First, I consider the Prosecutor’s incentives. I focus on the tradeoff she confronts and the implications of the informational asymmetry I model. Second, I turn to the Officer. I focus on the constraints he faces in pursuing his goals. Third,

⁵Both types warrant conviction on at least the low charge; there is no innocent type, so the model does not speak to wrongful conviction, though adding such a type is a natural extension.

⁶Costless alternative charging would strictly dominate $c = H$ whenever $\lambda > 0$; any realistic friction—the effort of parallel cases, or a hedge reading as the prosecution’s own doubt about the serious charge—restores the tradeoff.

I characterize equilibrium behavior. The analysis reveals three distinct types of behavior that can be supported. Finally, I present a variety of comparative statics that have implications for how we understand and evaluate patterns in prosecutorial choices.

4.1 The Prosecutor's Incentives

The Prosecutor faces two goals, which can potentially be in tension with each other. On one hand, the Prosecutor wants to secure convictions. The microfoundations for the Prosecutor's goals are both intrinsic and instrumental with respect to factors outside of the model, as described above. However, Prosecutors may also have an incentive to avoid over-convictions, which is the objective at the heart of the model. Some prosecutors have a distaste for over-charging individuals whose culpability does not warrant serious charges or who would be better off with less severe treatment in the criminal justice system. This interest is measured by the parameter λ . Thus, even setting aside the Prosecutor's informational disadvantage, she does not have a generically strictly dominant strategy. Her preferred strategy will depend on the relative weight of her electoral incentives for conviction relative to her taste for avoiding over-charging suspects.

That trade-off is compounded by the informational challenge the Prosecutor faces. The Prosecutor wants to know the suspect's type—because of λ —but anticipates that the Officer might want to mislead her. However, even if she believes that the Officer has not truthfully signaled the suspect's type, she still has some incentive to adopt the Officer's recommended charges, in order to cooperate with the Officer. This tension is at the core of the model's motivation, which is to understand how a prosecutor's priorities can undermine her relationship with the police.

Thus, if the Prosecutor believes with certainty the suspect is a high type (i.e., $\theta = H$), her choice is relatively simple. She is strictly better off choosing high-type charges (i.e., $c = H$), assuming the Officer has made high-type allegations.⁷ She cannot suffer a trust

⁷As we will see in the equilibrium analysis, low-type allegations against a high-type defendant only occur

cost from downgrading charges, and she does not have to worry about over-charging the suspect. However, when she believes that the suspect is a low type (i.e., $\theta = L$), her calculus is complicated by her competing interests. Given a sufficient interest in matching low-type suspects to low-type charges, she will be willing to forego the (potential) electoral benefit of securing convictions on high-type charges.

Remark 1 *The Prosecutor’s strategy will depend on both the value she places on matching low-type suspects to low-type charges and her belief about the suspect’s type.*

When we turn to equilibrium analysis, the question will be how the Prosecutor and the Officer can cooperate and credibly communicate information about the suspect’s type to coordinate charging decisions.

4.2 The Officer’s Incentives

The Officer’s latent incentives are somewhat simpler. The Officer values aggressive law enforcement and, to the extent measured by β , wants the Prosecutor to pursue high-type charges against suspects. But, that objective is complicated by three forces. First, the Officer does not like to see suspects acquitted (as measured by ϕ) and so has an incentive to avoid inducing the Prosecutor to pursue high-type charges if he thinks the suspect will not be convicted. That occurs with probability π when the suspect is a low-type. Thus, this consideration plays into the Officer’s decision calculus most clearly when the suspect is a low-type, because high-types can be convicted of any type of charge.

Second, because the Officer values matching charges to the suspect’s type, his taste for high-type charges by the Prosecutor is potentially in tension with his interest in aligning his own allegations with the suspect’s type—a comparison of the parameters β and γ . When $\beta < 0$, most of the Officer’s decisions become trivial, as do most of the model’s core tensions. The parameter γ measures how much the Officer cares about truthfully representing the

in equilibria where the Prosecutor’s strategy is non-responsive to the Officer’s strategy.

suspect's type and having the Prosecutor adopt the matching charge. Thus, the Officer faces a coordination problem with the Prosecutor. As we will see in the equilibrium analysis, their coordination dilemma is complicated by both their potentially conflicting tastes and the informational asymmetry about the suspect's type.

The third complication for the Officer is that he (like the Prosecutor) values trust and wants the Prosecutor to believe the allegations the Officer makes. This complicates his incentive to pursue high-type charges, because he wants to make allegations that the Prosecutor will adopt.

Taken together, for any given suspect, the Officer knows which charges he would like to recommend. He always wants the Prosecutor to follow his recommendation. Note, these complications of the Officer's taste for high-type charges all imply that he will prefer high-type charges for a weakly larger range of the parameter space than will the Prosecutor. However, there will be no cases in which the Officer prefers low-type charges but the Prosecutor does not. This is easy to verify via Table 1. There, we see that for an arbitrarily large β and an arbitrarily large λ , the Officer prefers high-type charges for a low-type suspect, whereas the Prosecutor would prefer low-type charges. However, under any condition where the Officer prefers high-type charges for a high-type suspect, so too would the Prosecutor. And this is the source of the most significant challenge the Officer faces. If the suspect is a low-type, he may not share the Prosecutor's preference for which charges to pursue in the case.

Remark 2 *The Officer may disagree with the Prosecutor about which charges to pursue in the event of a low-type suspect.*

Because his preferences for any given suspect might be in tension with those of the Prosecutor, he must contend with the possibility that the Prosecutor will not trust his allegations. This is the matter I study in the equilibrium analysis that follows.

4.3 Equilibrium

In characterizing equilibrium behavior, I focus on perfect Bayesian equilibria. It will be convenient to focus on two key parameters of substantive interest—the Prosecutor’s interest in matching low-type suspects to low-type charges (λ) and the courts’ aversion to convicting low-type suspects of high-type crimes (π). Three kinds of behavior can happen in equilibrium: the Officer’s allegations can reveal his private information fully (separation), partially (semi-separation), or not at all (pooling and other uninformative behavior). Panel (a) of Figure 1 summarizes the most informative equilibrium at each point of the parameter space, relative to the two parameters of interest, λ and π . Panel (b) overlays the same space with the regions in which the model’s uninformative (babbling) equilibria exist; as discussed below (Proposition 4), these coexist with the equilibria shown in panel (a) rather than replacing them. Table 2 summarizes the equilibria and the conditions under which each can be sustained. The Officer’s taste for severity (β) and taste for justice (γ) jointly determine the threshold $\bar{\pi}$ that separates the fully informative region of the parameter space from the rest, while a second threshold in λ , derived below, separates pooling from partial communication.

Throughout, I restrict attention to equilibria in which the Officer’s allegations carry their natural meaning: when the Officer’s strategy is informative, a high-type allegation indicates a high-type suspect, and I set aside equilibria sustained only by the Officer’s exact payoff indifference between messages.⁸

The first thing to note is that a separating equilibrium requires that the courts be sufficiently averse to convicting low-type defendants of high-type charges. It is the threat of a court declining to convict a low-type defendant on serious charges that mitigates both the

⁸The restriction does substantive work in one region of the parameter space. When the Officer’s aversion to acquittal is extreme relative to his taste for severity and the trust cost ($\pi\phi \geq \beta + \delta$, along with $\eta \geq \epsilon$ and $\lambda \geq (1 - \pi)\eta + \delta$), the model admits an *inverted* separating equilibrium in which the Officer alleges $a = L$ against high types and $a = H$ against low types, and the Prosecutor charges against the allegation: a false high-type allegation is so self-punishing in that regime that it credibly signals a low type. This equilibrium is strict and survives belief-based refinements, but it requires the police to systematically understate the cases they care most about and overstate the ones they do not, an inversion of institutional language without an empirical analogue in the police-prosecutor setting; equivalently, it can be eliminated by assuming $\pi\phi < \beta + \delta$.

Equilibrium	Conditions	Officer's strategy	Prosecutor's strategy	Outcome for low-type suspect
Separating (Prop. 1)	$\pi \geq \bar{\pi};$ $\lambda > \underline{\lambda}$	$a = H$ if $\theta = H;$ $a = L$ if $\theta = L$	$c = a$ (follows Officer)	Charged L ; always convicted
Pooling (Prop. 2)	$\pi < \bar{\pi};$ $\lambda < \lambda^*$	$a = H$ always	$c = H$ always	Charged H ; convicted w.p. $1 - \pi$
Semi-separating (Prop. 3)	$\pi < \bar{\pi};$ $\lambda \geq \lambda^*$	$a = H$ if $\theta = H;$ $a = H$ w.p. σ_O^* if $\theta = L$	$c = L$ if $a = L;$ $c = H$ w.p. σ_P^* if $a = H$	If $a = L$: charged L , convicted. If $a = H$: charged H w.p. σ_P^* ; convicted w.p. $1 - \pi$
Uninformative (Prop. 4)	All- H form: $\lambda \leq \lambda^*;$ all- L form: λ large (see Prop. 4)	a independent of θ	c independent of a ; $c = H$ (all- H form); $c = L$ (all- L form)	Charge depends on prior p , not on θ

Notes: $\bar{\pi} \equiv \frac{\beta - \gamma}{\phi}$; $\underline{\lambda} \equiv (1 - \pi)\eta - \epsilon$; $\lambda^* \equiv \frac{p\eta + (1-p)(1-\pi)\eta + \delta}{1-p}$; $\bar{p} = \frac{\lambda - (1-\pi)\eta - \delta}{\lambda + \pi\eta}$, with $p \geq \bar{p} \Leftrightarrow \lambda \leq \lambda^*$. In the semi-separating equilibrium, $\sigma_O^* = \frac{p(\eta + \delta)}{(1-p)(\lambda - (1-\pi)\eta - \delta)}$ and $\sigma_P^* = \frac{\gamma + \delta}{\beta + \delta - \pi\phi}$; at $\lambda = \lambda^*$, $\sigma_O^* = 1$ and the Officer's allegation strategy coincides with pooling. Note that $\lambda^* > \underline{\lambda}$ for all parameter values. The separating equilibrium requires $\lambda > \underline{\lambda}$; for $\pi \geq \bar{\pi}$ and $\lambda < \underline{\lambda}$, behavior pools on high charges (the triangular region in Figure 1). The all-low uninformative equilibrium is never unique: wherever it exists it coexists with the pooling, semi-separating, or separating equilibrium, and where it coexists with separation it fails the Intuitive Criterion (Remark 4).

Table 2: Summary of equilibrium behavior.




Figure 1: *Equilibrium space in the baseline model.* Panel (a) shows the most informative equilibrium at each point of the parameter space, for the focal case $\beta > \gamma > 0$. To the left of $\bar{\pi}$, the boundary between pooling and semi-separation is $\lambda^*(\pi)$; to the right of $\bar{\pi}$, separation is possible whenever $\lambda > \underline{\lambda}(\pi)$, and behavior pools on high charges in the triangular region below. The two boundary lines are parallel (each has slope $-\eta$ in π), with $\lambda^* - \underline{\lambda} = \frac{p\eta + \delta}{1-p} + \epsilon > 0$. The points marked λ^* and $\underline{\lambda}$ show those values evaluated at $\pi = 0$; the figure assumes $\eta > \epsilon$, so that $\underline{\lambda} > 0$ at $\pi = 0$. Panel (b) retains the same regions and overlays where the two uninformative equilibria exist (Proposition 4): the all-high form ($a = H, c = H$), hatched “/”, coincides with the pooling region ($\pi < \bar{\pi}, \lambda \leq \lambda^*$), the same all-high charging outcome that also arises in the triangular region ($\pi \geq \bar{\pi}, \lambda < \underline{\lambda}$) where the Prosecutor escalates a known low type; the all-low “standing down” form ($a = L, c = L$), hatched “\”, exists above a threshold parallel to and below λ^* (dash-dotted), spanning all π . Where the hatching crosses, both forms coexist. Each uninformative equilibrium coexists with the most informative equilibrium shown in panel (a); where the all-low form coexists with separation it fails the Intuitive Criterion (Remark 4).

Officer's and Prosecutor's interest in securing an over-conviction against a low-type suspect. When the court is sufficiently likely to balk at such a conviction, the Officer no longer has an incentive to try to mislead the Prosecutor. Moreover, not only is this judicial disposition a sufficient condition for separation, it is also necessary. Under no conditions can separation be supported when the court is more willing to convict low types of serious charges. It will be useful to define that threshold, $\bar{\pi}$, which is increasing in the Officer's taste for severity (β); as that taste grows, a court more strongly averse to over-conviction is necessary to make the Officer willing to honestly communicate the suspect's type.

Definition 1 *Define $\bar{\pi} \equiv \frac{\beta-\gamma}{\phi}$. This is the critical value of π , such that separation is never possible when $\pi < \bar{\pi}$.*

The first result follows: when the court is sufficiently averse to over-conviction and the Prosecutor sufficiently values avoiding over-charges, behavior is perfectly separating. Because the Officer's allegations then identify the suspect's type with certainty, the Prosecutor optimally adopts them.

Proposition 1 *If the courts are sufficiently averse to convicting low-type defendants of high-type charges ($\pi \geq \bar{\pi}$) and the Prosecutor sufficiently values avoiding over-charges ($\lambda > \underline{\lambda}$), there exists a perfect Bayesian equilibrium with pure separation. The Officer matches his allegations to the suspect's true type, and the Prosecutor matches her charges to the Officer's allegations. Because only high-types face high-type charges, and all low-types face low-type charges, the suspect is always convicted.*

Importantly, though, for a sufficiently low λ , separation requires an even stronger judicial aversion to over-conviction. As we can see in the triangular area to the right of $\bar{\pi}$ at the bottom of panel (a) of Figure 1, separation also requires that the Prosecutor be sufficiently interested in ensuring that a low-type suspect is charged with low-type charges (λ). The judicial aversion to over-conviction that is necessary to support separation is greater for lower values of λ . That is, increasing the extent to which the Prosecutor values matching low-type suspects to low-type charges weakens the judicial disposition necessary to support separation until a critical level is reached, at which point further increases in λ do not affect

the condition on the courts.

Remark 3 *Increasing the Prosecutor’s preference for not over-charging the suspect weakly increases the range of judicial dispositions that can support a separating equilibrium in which the Officer files accurate allegations against the suspect and the Prosecutor adopts the charges the Officer makes.*

It also bears noting that whenever the Officer leans against high-type charges (i.e., $\beta < 0$), then separation can be supported for any judicial disposition. That is, $\pi > \bar{\pi}$ for any $\beta < 0$. Thus, when considering other regions of the parameter space below, the most interesting dynamics will involve cases where $\beta > 0$.

Proposition 1 reveals another dynamic, which is that for a sufficiently pro-conviction court (i.e., a low enough π), some types of Prosecutors are willing to charge a suspect with high-type charges even when she is *certain* that the suspect is a low-type. This requires that the Prosecutor not value matching low-type charges to low-type suspects enough, relative to the court’s willingness to convict—i.e., that λ is sufficiently small relative to π .

Corollary 1 *For a sufficiently small λ and π , the Prosecutor is willing to pursue high-type charges against a suspect she knows to be a low-type.*

It is worth pausing to note why there is no pooling equilibrium in which the Officer always chooses $a = H$ and the Prosecutor always responds with $c = L$. The intuition is straightforward: if the Prosecutor were going to charge $c = L$ regardless, the Officer would strictly prefer to choose $a = L$ as well, because choosing $a = H$ only to have the Prosecutor downgrade the charges imposes the trust cost $\delta > 0$ on the Officer. Thus, pooling on $a = H$ with $c = L$ unravels from the Officer’s side. The only pooling equilibrium in which the Officer chooses $a = H$ is one in which the Prosecutor follows suit with $c = H$, because it is only in that case that the Officer’s high-type allegations are rewarded rather than penalized.

Corollary 1 gives a sufficient condition for pooling, but pooling can arise under weaker

conditions. It is useful to define the critical value of λ below which the Prosecutor prefers to gamble on a high-type conviction *even against a suspect she knows to be a low type*.

Definition 2 Define $\underline{\lambda} \equiv (1 - \pi)\eta - \epsilon$. Whenever $\lambda < \underline{\lambda}$, the Prosecutor prefers high-type charges even when certain that $\theta = L$.

Pooling does not, however, require so extreme a preference. When the Officer pools on high-type allegations, the allegation carries no information, and the Prosecutor's posterior belief equals her prior, p . What sustains pooling on the Prosecutor's side is therefore that she prefer high-type charges *given her prior*—a weaker requirement than the certainty benchmark of Corollary 1, because with probability p the suspect is in fact a high type whose serious conviction is assured. The critical value of λ for this comparison plays a central organizing role in the analysis, so I define it explicitly.

Definition 3 Define $\lambda^* \equiv \frac{p\eta + (1-p)(1-\pi)\eta + \delta}{1-p}$. Given an uninformative high-type allegation, the Prosecutor prefers high-type charges if and only if $\lambda \leq \lambda^*$. Note that $\lambda^* > \underline{\lambda}$ for all parameter values, and that λ^* is increasing in η , δ , and p and decreasing in π .

The intuition behind the pooling equilibrium is straight-forward. When π is small enough, the Officer is willing to gamble on having the suspect prosecuted with high-type charges, even if the suspect is a low type. The high-type suspect will be convicted with certainty, while the small π justifies the risk of acquitting a low-type. Formally, the critical value of π for the Officer is $\frac{\beta - \gamma}{\phi} \equiv \bar{\pi}$. So, as the Officer's taste for severity grows, he is willing to take that chance for a wider range of π —less likely convictions of low-type suspects on high-type charges. Naturally, as the Officer becomes more interested in matching charges to suspect types (γ), we find the opposite relationship. In order for the Officer to be willing to pool, it must also be the case that the Prosecutor be willing to choose high-type charges, understanding that the Officer is pooling—that is, $\lambda \leq \lambda^*$. As the value of securing high-type convictions (η) increases, or as the trust cost of downgrading allegations (δ) increases, the threshold λ^* rises and the Prosecutor is willing to pool for a wider range of leniency

preferences.

Proposition 2 *For $\pi < \bar{\pi}$ and $\lambda < \lambda^*$, there exists a pooling equilibrium in which the Officer always chooses $a = H$ and the Prosecutor always chooses $c = H$. Low-type suspects are probabilistically convicted of high-type crimes and high-type suspects are always convicted.⁹*

What happens when the Prosecutor's leniency interest exceeds λ^* ? Separation remains impossible for $\pi < \bar{\pi}$, and pooling on high charges cannot be sustained, because a Prosecutor facing an uninformative high-type allegation would deviate to low-type charges. In this situation, we have a hybrid, semi-separating equilibrium. The Officer always makes high-type allegations ($a = H$) against a high-type suspect ($\theta = H$) and probabilistically makes high-type allegations against a low-type suspect ($\theta = L$), mixing at exactly the rate that leaves the Prosecutor indifferent after a high-type allegation. The result is that if the Prosecutor sees low-type allegations ($a = L$), then she knows the suspect is a low-type, and she will always choose low-type charges ($c = L$). However, if she sees high-type allegations, she knows there is some chance that the suspect is in fact a low type. Thus, she probabilistically drops charges from high-type charges to low-type charges, mixing at exactly the rate that leaves the Officer facing a low-type suspect indifferent. Because some low-type suspects will be charged as high types, there is positive probability that a low-type suspect is convicted of high-type charges.

Proposition 3 *For $\pi < \bar{\pi}$ and $\lambda \geq \lambda^*$, there exists a semi-separating equilibrium. In this equilibrium, the Officer always chooses $a = H$ when $\theta = H$. When $\theta = L$, the Officer chooses $a = H$ with probability $\sigma_O^* = \frac{p(\eta+\delta)}{(1-p)(\lambda-(1-\pi)\eta-\delta)}$. Upon observing $a = L$, the Prosecutor believes with certainty that $\theta = L$ and chooses $c = L$. Upon observing $a = H$, the Prosecutor believes that $\theta = H$ with probability $\hat{p} = \frac{p}{p+(1-p)\sigma_O^*}$ and chooses $c = H$ with probability $\sigma_P^* = \frac{\gamma+\delta}{\beta+\delta-\pi\phi}$. At $\lambda = \lambda^*$, $\sigma_O^* = 1$, so the Officer's allegation strategy coincides with pooling on high-type allegations (Proposition 2), though the Prosecutor continues to mix at σ_P^* ; as λ grows beyond λ^* , σ_O^* strictly declines.*

The boundary behavior described in the last sentence of the proposition is worth em-

⁹For $\pi \geq \bar{\pi}$, pooling on high charges can also be sustained when $\lambda < \underline{\lambda}$, because the Prosecutor would escalate even a revealing low-type allegation (Corollary 1); this is the triangular region of panel (a) of Figure 1. For $\pi \geq \bar{\pi}$ and $\lambda > \underline{\lambda}$, pooling can be sustained only by off-path beliefs that fail the Intuitive Criterion, and I set those equilibria aside.

phasizing, because it organizes the left-hand side of panel (a) of Figure 1. The thresholds for pooling and semi-separation are the same threshold: λ^* is simultaneously the largest leniency interest consistent with the Prosecutor rubber-stamping uninformative high-type allegations and the smallest leniency interest consistent with the Officer's mixing probability being a proper probability. Below λ^* , the Officer misrepresents every low-type suspect and the Prosecutor goes along; above λ^* , the Officer misrepresents only some low-type suspects, and the rate of misrepresentation σ_O^* falls continuously as λ rises. In the limit of a very large λ , the Officer almost always reports low-type suspects truthfully—but the Prosecutor continues to discount high-type allegations, following them only with probability σ_P^* . A prosecutor's leniency commitment thus disciplines police misrepresentation rather than foreclosing communication.

Two features of the equilibrium mixing probabilities deserve note, because they drive comparative statics developed below. First, σ_O^* is increasing in δ : the more costly it is for the Prosecutor to downgrade allegations, the more license the Officer has to misrepresent, because the Prosecutor's trust cost makes her more reluctant to act on her skepticism. Second, σ_P^* is also increasing in δ (for $\pi < \bar{\pi}$): the Prosecutor follows high-type allegations more often when downgrading them is costlier. Both effects work in the same direction. Higher trust costs mean more misrepresentation and more acquiescence; I return to this in the next section.

The model also admits equilibria in which the Officer's allegation carries no information at all—babbling equilibria, familiar from the cheap-talk literature. In such an equilibrium the Prosecutor's charge depends only on her prior. The all-high form—in which she charges $c = H$, requiring a prior high enough, $p \geq \bar{p}$, a condition equivalent to $\lambda \leq \lambda^*$ —is simply the pooling equilibrium of Proposition 2 viewed through an informational lens. The all-low form, in which she charges $c = L$, is substantively distinct.

Proposition 4 *Define $\bar{p} = \frac{\lambda - (1 - \pi)\eta - \delta}{\lambda + \pi\eta}$, and note $p \geq \bar{p} \Leftrightarrow \lambda \leq \lambda^*$. An uninformative equi-*

librium, in which the Officer's allegation is independent of θ and the Prosecutor's charge depends only on her prior, exists in at least one of two forms in which charges match allegations. (i) If $\pi < \bar{\pi}$ and $p \geq \bar{p}$ (equivalently $\lambda \leq \lambda^*$), the Officer pools on high-type allegations and the Prosecutor always charges $c = H$ —the pooling equilibrium of Proposition 2. (ii) If $\lambda \geq \max \left\{ (1 - \pi)\eta + \delta, \lambda^* - \frac{\delta + \epsilon}{1 - p} \right\}$ —a condition satisfied, in particular, whenever $\lambda \geq \lambda^*$ —there exists an equilibrium in which the Officer always alleges $a = L$ and the Prosecutor always charges $c = L$. The all-high form is the pooling equilibrium of Proposition 2, not a distinct equilibrium; the all-low form is never unique: wherever it exists, it coexists with the pooling, semi-separating, or separating equilibrium of Propositions 1–3.

The all-low form of the uninformative equilibrium has a distinctive logic. The Prosecutor's commitment to leniency (or her low prior) means that no allegation can move her to serious charges; anticipating that high-type allegations will be downgraded—at trust cost δ to both players—the Officer pre-emptively under-alleges every suspect, including the genuinely high types. Communication breaks down because the Prosecutor's posture renders the Officer's information useless to him. High-type suspects are systematically under-charged in this equilibrium, an outcome that resembles the “standing down” behavior described in the introduction: the police continue to make arrests but stop advancing serious allegations they expect the prosecutor to repudiate.

A selection question therefore arises where equilibria coexist. Where the uninformative equilibrium coexists with separation ($\pi > \bar{\pi}$), the Intuitive Criterion eliminates it; where it coexists with semi-separation ($\pi < \bar{\pi}$ and $\lambda \geq \lambda^*$), standard belief-based refinements have no bite (the formal arguments are in the proof of Proposition 4). In what follows, I focus on the most informative equilibrium at each point of the parameter space, following the standard convention in communication games; but the coexistence itself is substantively meaningful, a point developed in the next section.

Remark 4 *Among equilibria in which allegations carry their natural meaning, only the separating equilibrium survives the Intuitive Criterion when $\pi > \bar{\pi}$ and $\lambda > \underline{\lambda}$. When $\pi < \bar{\pi}$ and $\lambda \geq \lambda^*$, the semi-separating and uninformative equilibria coexist and standard refinements do not select between them; the analysis below focuses on the most informative equilibrium.*

4.4 Comparative Statics and Empirical Implications

One of the model’s most striking features is that three equilibria are observationally equivalent under perfect alignment. In the separating, all-high pooling, and all-low equilibria alike, the Prosecutor’s charge always matches the Officer’s allegation ($c^* = a^*$). The absence of downgrades therefore cannot by itself distinguish an informative equilibrium—in which the Officer truthfully conveys the suspect’s type and the Prosecutor adopts it—from a non-informative one, in which the Officer either over-alleges every low type (all-high pooling) or conveys nothing (all-low). Only the semi-separating equilibrium produces mismatch, so it is observed disagreement between police and prosecutors, rather than the alignment we usually see, that is empirically informative.

Result 1 *Perfect alignment between police allegations and prosecutors’ charges does not discriminate between informative and non-informative communication: the separating, all-high pooling, and all-low equilibria all satisfy $c^* = a^*$. Mismatch between an allegation and the charge arises only in the semi-separating equilibrium, so it is observed disagreement, not alignment, that carries information about the underlying regime.*

Propositions 1 through 4 describe the range of equilibrium behavior the model supports, but not all of it arises for every parameter value. The first comparative static concerns the Officer’s taste for severity: because separation requires $\pi > \bar{\pi} = \frac{\beta - \gamma}{\phi}$, a weaker taste for severity (lower β) makes separation easier to sustain, and once $-\gamma < \beta < \gamma$ it holds for any court (Corollaries 3–4). As β rises, separation requires an increasingly averse court.

Result 2 *As an Officer’s taste for severity weakens, or the courts become more averse to convicting low-type suspects of high-type charges, separating behavior is more likely in which the Officer always matches his allegations to the suspect type and the Prosecutor always adopts the Officer’s allegations.*

With respect to panel (a) of Figure 1, this result concerns the location of the vertical

line that partitions the space. As β increases, that line moves to the right. Notably, the downward-sloping boundary that describes the space to the left of the vertical line does not depend on β or γ . Thus, those parameters simply affect whether behavior is captured by the fully informative equilibrium behavior or the less informative behavior in pooling or semi-separating equilibria.

Turning to the left-hand side of panel (a) of Figure 1, the central parameter is the Prosecutor's interest in avoiding over-charges, λ , which operates through the threshold λ^* described above. Its effect is not merely to shift charging outcomes: a commitment to avoiding over-charges disciplines the quality of the information the police transmit. The countervailing cost is borne by the trust relationship and by high-type enforcement. Above λ^* , the Prosecutor sometimes downgrades high-type allegations (with probability $1 - \sigma_P^*$), incurring the trust cost δ on path, and the coexisting uninformative equilibrium in which police preemptively under-allege every suspect becomes available.

Result 3 *For a court below the Officer's threshold ($\pi < \bar{\pi}$), the Prosecutor's interest in avoiding over-charges affects communication through the threshold λ^* : for $\lambda < \lambda^*$ allegations are uninformative and all suspects face high-type charges, and for $\lambda \geq \lambda^*$ allegations are partially informative, with the Officer's misrepresentation of low-type suspects strictly declining as λ increases. For $\pi \geq \bar{\pi}$ the operative boundary is instead $\underline{\lambda}$, separating pooling on high charges from separation.*

Comparative statics on the location of the threshold λ^* yield a series of important implications. Recall from Definition 3 that λ^* is increasing in the electoral value of high-type convictions (η), increasing in the trust cost (δ), increasing in the prior (p), and decreasing in the court's aversion to over-conviction (π). Consider first the electoral incentive. Increasing η raises λ^* and so expands the pooling region (panel (a) of Figure 2): a prosecutor with stronger electoral incentives to secure serious convictions requires a stronger leniency commitment before she will stop rubber-stamping high-type allegations. Moreover, the marginal effect of η on λ^* is larger when π is small. The more willing the courts are to convict low types of serious charges, the greater the impact of electoral incentives on prosecutorial

punitiveness.

Result 4 *Increasing the electoral returns to convicting criminal suspects on high-type charges expands the range of prosecutorial preferences for which Officers pool on high-type allegations and Prosecutors pool on high-type charges. That effect is exacerbated where courts are more willing to convict low-type suspects of high-type charges.*




Figure 2: *Electoral incentives and the communication threshold.* Panel (a) plots the semi-separation/pooling boundary λ^* against the Prosecutor's electoral benefit η , holding p and δ fixed. Because λ^* is affine in η (Definition 3), the three curves—drawn for a court below, at, and above the Officer's threshold $\bar{\pi}$ —are straight lines sharing the vertical intercept $\frac{\delta}{1-p}$ and fanning upward; the slope $\frac{1-(1-p)\pi}{1-p}$ is flatter for a court more averse to over-conviction (larger π). For a court below the Officer's threshold ($\pi < \bar{\pi}$), λ^* is the boundary between pooling (below) and semi-separation (above); the $\pi = \bar{\pi}$ and $\pi > \bar{\pi}$ lines are drawn to display the comparative static, but for those courts the operative boundary for the most informative equilibrium is $\underline{\lambda}$ (separation above, pooling on high charges below), not λ^* . Panel (b) plots the prior threshold $\bar{p}$ above which the Prosecutor pools on high-type charges, as a function of η for fixed λ ; it is decreasing in η , so the shaded pooling region ($p \geq \bar{p}$, equivalently $\lambda \leq \lambda^*$) expands as electoral stakes rise. A single $\bar{p}(\eta)$ curve is shown because it is decreasing for every court, and the equivalence $p \geq \bar{p} \Leftrightarrow \lambda \leq \lambda^*$ again characterizes the pooling boundary for $\pi < \bar{\pi}$.

Figure 2 collects these electoral-incentive comparative statics and extends them to the prior. Panel (b) plots the mirror-image object, the prior threshold $\bar{p}$ above which the Prosecutor pools on high-type charges. Because $\frac{\partial \bar{p}}{\partial \eta} = \frac{\pi\delta - \lambda}{(\lambda + \pi\eta)^2} < 0$ throughout the informative

region, $\bar{p}$ falls as η grows, so pooling on serious charges is sustained across a wider range of prosecutorial priors. Panels (a) and (b) are two views of the same force: because $p \geq \bar{p}$ if and only if $\lambda \leq \lambda^*$, a rise in η that lifts the λ^* boundary also lowers the prior required to sustain pooling. The separating region contracts as well: the escalation threshold $\underline{\lambda} = (1 - \pi)\eta - \epsilon$ is likewise increasing in η , so both boundaries of the informative region shift upward as electoral stakes rise, and separating and semi-separating equilibria alike become harder to sustain.

The electoral incentive also reshapes behavior *within* the equilibria, in ways that are less immediate. In the semi-separating equilibrium the Officer’s misrepresentation rate rises with η : $\frac{\partial \sigma_O^*}{\partial \eta} = \frac{p(\lambda - \pi\delta)}{(1-p)(\lambda - (1-\pi)\eta - \delta)^2} > 0$ on the informative domain, so a larger electoral benefit leads to more misrepresentation of low-type suspects at the margin, for any fixed (λ, π) that still supports semi-separation. The Prosecutor’s acquiescence rate, $\sigma_P^* = \frac{\gamma + \delta}{\beta + \delta - \pi\phi}$, by contrast, is independent of η : it is fixed by the Officer’s indifference condition, into which the electoral benefit does not enter. The sharpest effect is on the “standing down” equilibrium. Its existence requires $\lambda \geq \max \left\{ (1 - \pi)\eta + \delta, \lambda^* - \frac{\delta + \epsilon}{1 - p} \right\}$, and both bounds are increasing in η . A stronger electoral incentive thus demands an *even more* lenient prosecutor before police will pre-emptively under-allege: the stand-down equilibrium contracts as the electoral returns to serious convictions grow, and is therefore a phenomenon of low-electoral-stakes, high-leniency environments—an observation I return to when considering which combinations of η and λ are politically plausible.

Result 5 *Stronger electoral returns to high-type convictions increase the rate at which police misrepresent low-type suspects in the semi-separating equilibrium, but leave unchanged the rate at which prosecutors adopt high-type allegations; and they shrink the range of prosecutorial preferences over which the uninformative “standing down” equilibrium exists.*

The trust parameter δ produces a particularly non-obvious result. In particular, the model implies that the more both players value their working relationship, the *worse* they communicate. The trust cost enters the analysis in three places, and all three run against information transmission. First, λ^* is increasing in δ : a higher trust cost expands the pooling

region, because a Prosecutor who finds it costlier to repudiate police allegations requires a stronger leniency commitment before she will do so. Second, within the semi-separating equilibrium, the Officer's misrepresentation probability σ_O^* is increasing in δ : the Officer exploits the Prosecutor's reluctance to downgrade, misrepresenting low-type suspects more often. Third, the Prosecutor's acquiescence probability σ_P^* is also increasing in δ : she follows high-type allegations more often because this makes contradicting them costly. (Notably, this is true despite the allegations being less credible.), In other words, the institutional bond between police and prosecutors creates a cost associated with prosecutorial skepticism which the police have to account for, as opposed to being a conduit for information.

Result 6 *Increases in the cost of lost trust between police and prosecutors (δ) degrade the informativeness of police allegations: they expand the range of prosecutorial preferences for which communication is fully uninformative and, where communication is partially informative, increase both the rate at which police misrepresent low-type suspects and the rate at which prosecutors adopt high-type allegations.*

Finally, the analysis points to the importance of the Prosecutor's prior belief that the suspect is a high-type. The prior enters the equilibrium structure directly: λ^* is increasing in p , so prosecutors who expect the police to bring them mostly high-culpability suspects pool on serious charges for a wider range of leniency preferences; and within the semi-separating equilibrium, σ_O^* is increasing in p , so a more favorable prior also licenses more police misrepresentation at the margin. Of course, there are institutional factors that affect that belief, most evidently the quality of police work. But, there are selection mechanisms shaping who becomes a prosecutor in the first place that can shape those beliefs. We might also imagine that broader political forces play a role in framing the context that gives rise to a prosecutor's expectations about the suspects that police bring to them for prosecution. Evaluating the nature of the Prosecutor's beliefs and how they might be endogenously shaped by parts of the criminal justice and political-legal system is an important avenue for research.

Political plausibility of high-leniency, high-stakes environments. A natural objection to reading these comparative statics parameter by parameter is that λ and η are unlikely to vary independently. Both are shaped by the same political environment, and an electorate that places a high value on serious convictions (large η) is unlikely to return a strongly reform-minded prosecutor (large λ). The high- λ , high- η corner of the parameter space may therefore be nearly empty in practice. The model does not endogenize electoral selection, but the concern can be accommodated without re-solving the game by imposing a reduced-form feasibility frontier $\lambda(\eta) = \lambda_0 - k\eta$ with $k > 0$: an electorate that rewards serious convictions more heavily returns prosecutors with weaker leniency commitments. Informative communication requires $\lambda(\eta) \geq \lambda^*(\eta)$; because λ^* rises in η while the feasible λ falls, this holds only for

$$\eta \leq \hat{\eta} \equiv \frac{\lambda_0 - \frac{\delta}{1-p}}{k + \frac{1-(1-p)\pi}{1-p}}.$$

The standing-down equilibrium, by contrast, requires only $\lambda(\eta) \geq B_{SD}(\eta)$, where $B_{SD} = \max\{(1 - \pi)\eta + \delta, \lambda^* - \frac{\delta + \epsilon}{1-p}\} < \lambda^*$; it therefore survives to a strictly higher cutoff $\hat{\eta}_{SD} > \hat{\eta}$. Figure 3 overlays the feasibility frontier on both thresholds. Informative communication becomes infeasible first, at $\hat{\eta}$, and standing down somewhat later, at $\hat{\eta}_{SD}$; beyond that, the only feasible prosecutor is too punitive to sustain either, and police and prosecutors pool on serious charges. Both informative-leniency and standing-down equilibria are thus phenomena of high- λ , low- η environments. Fully endogenizing selection—deriving $\lambda(\eta)$ from voter preferences over charging outcomes—is a natural but separate undertaking; the linear frontier here is meant only to locate the feasible region.

5 Discussion and Conclusion

In recent years, political scientists have focused renewed attention on the politics of the criminal justice system. Alternatively examining policing and prosecution, that work highlights the ways in which discretion by powerful actors in American civic life affects the application




Figure 3: *Feasibility of informative and standing-down equilibria under political selection.* The solid line is the informative-communication threshold $\lambda^*(\eta)$; the dash-dotted line is the standing-down threshold $B_{SD}(\eta) = \max\{(1 - \pi)\eta + \delta, \lambda^*(\eta) - \frac{\delta + \epsilon}{1 - p}\}$, which lies below λ^* ; the falling dashed line is the reduced-form feasibility frontier $\lambda(\eta) = \lambda_0 - k\eta$, capturing the idea that electorates valuing serious convictions more heavily return prosecutors with weaker leniency commitments. Informative communication requires $\lambda(\eta) \geq \lambda^*(\eta)$, feasible for $\eta \leq \hat{\eta}$; the standing-down equilibrium requires only $\lambda(\eta) \geq B_{SD}(\eta)$, feasible for the wider range $\eta \leq \hat{\eta}_{SD}$. Standing down therefore survives past the point at which informative communication becomes infeasible, and both give way to pooling on serious charges once $\eta > \hat{\eta}_{SD}$. Both are phenomena of high- λ , low- η environments.

of law. The analysis here helps put some of those insights together to illustrate how the interaction of these two critical actors in the American legal system shapes the way the law is applied to individuals. The results have implications for a number of salient subjects of study in the field.

Empirical implications. While the model's most immediate contribution to empirical research is a cautionary one—Result 1 demonstrates that alignment between police and prosecutors is consistent with very different underlying dynamics—the analysis also yields testable predictions. Most directly, the model predicts a threshold relationship between a prosecutor's leniency preferences and the informativeness of police-prosecutor communica-

tion. Below the threshold, jurisdictions should exhibit lockstep alignment: police allege and prosecutors charge serious crimes regardless of underlying culpability. Above it, communication is partially informative, and the model identifies a distinctive observable signature: prosecutors sometimes downgrade police allegations. Because mismatch between allegations and charges occurs only in the semi-separating equilibrium (within the model’s natural-language equilibria), observed downgrade rates are direct evidence of partially informative communication. Perfect alignment, though is difficult to interpret theoretically. Related, and counter-intuitively, conditional on being in the communicative regime, police misrepresentation declines as the prosecutor’s leniency commitment grows, so more strongly reform-oriented prosecutors should receive *more* truthful allegation behavior from police, even as they more often act against police allegations. Third, the model predicts that stronger police-prosecutor trust relationships cut against information: where the political cost of public divergence between the institutions is high, lockstep alignment is more likely, and any communication that survives is less informative. Finally, the courts’ disposition tendencies interact with prosecutorial preferences: where courts are sufficiently averse to convicting low-culpability defendants of serious charges, fully truthful communication is possible and charging patterns track culpability. Cross-jurisdictional variation in adjudication (for example, brought about by differences in public defender resources or judicial caseloads) provides a natural source of leverage for testing these predictions.

Possibilities for reform prosecutors. To be sure, not only do police potentially adjust their behavior in anticipation of prosecutorial discretion, but so too do prosecutors respond to expectations about the electoral landscape (e.g., Gordon and Huber 2002, 2009). The epidemic of mass incarceration in the US has motivated some candidates for prosecutorial offices to run on reform platforms, promising to avoid over-punishing small crimes. At the same time, punitive approaches to law enforcement, especially in the wake of high-profile events or spikes in crime, remain effective in American electoral politics. And, police make

use of public support for crime-fighting to build political capital for their own policy goals (e.g., Cheng 2023). Those competing tensions can combine to shape what is possible for advocates for criminal justice reform.

While the analysis in this paper does not directly speak to those debates, several of the findings have implications for how they might play out. For example, the model demonstrates that cooperation between the police and prosecutors—in the sense that police can credibly communicate their street-level information about criminal suspects to prosecutors—depends on the disposition tendencies of the courts. One of the more challenging insights, for example, comes from Result 1, which demonstrates the observational equivalence of a number of theoretically distinct kinds of behavior. There, we saw that observed alignment between prosecutor and police behavior alone does not discriminate among informative and non-informative equilibria. Thus, if we want to understand the nature of cooperation and communication between police and prosecutors, empirical scholars and reformers will need to investigate patterns that incorporate broader features of the criminal justice system. The implication is that debates about law enforcement reform or prosecutorial reform cannot take place in isolation, nor can they take place independent of considerations about the workload and resources that courts have. While this point complicates matters, it underscores the extent to which police, prosecutors, and courts are intimately linked as components of the political-legal bureaucracy in America—a system whose internal information and incentive problems are precisely the domain of public-choice analyses of bureaucracy (e.g., Adamson and Rentschler 2023, Weingast 1984).

Directions for future research. The analysis suggests several extensions. First, because the officer’s private culpability assessment is exactly what a prosecutor needs to tailor charges, the model speaks to the police role in setting the agenda for plea bargaining; the communication problem studied here is logically prior to, and complementary to, existing formalizations of bargaining and pleas (e.g., Agan, Doleac and Harvey 2023, Little and Simpson

2024). Second, the model takes police effort and the arrest pool as given, but expectations about prosecutorial priorities plausibly shape both; extending the logic to endogenous police effort would connect it to the literature on bureaucratic expertise and discretion (e.g., Gailmard and Patty 2007, Mummolo 2018).¹⁰ Third, the officer’s information is taken as given; letting him choose what evidence to gather and disclose would recast the problem as one of Bayesian persuasion (e.g., Kamenica and Gentzkow 2011), a natural complement for studying how reforms such as body cameras or mandatory reporting reshape police-prosecutor dynamics.

¹⁰Recall that Result 2 already shows that officer preferences govern cooperation.

References

- Adamson, Jordan and Lucas Rentschler. 2023. “Criminal justice from a public choice perspective: An introduction to the special issue.” *Public Choice* 196(3–4):223–227.
- Agan, Amanda, Jennifer L. Doleac and Anna Harvey. 2023. “Misdemeanor prosecution.” *The Quarterly Journal of Economics* 138(3):1453–1505.
- Ba, Bocar A., Dean Knox, Jonathan Mummolo and Roman Rivera. 2021. “The role of officer race and gender in police-civilian interactions in Chicago.” *Science* 371(6530):696–702.
- Baker, Scott and Claudio Mezzetti. 2001. “Prosecutorial resources, plea bargaining, and the decision to go to trial.” *Journal of Law, Economics, and Organization* 17(1):149–167.
- Bar-Gill, Oren and Oren Gazal Ayal. 2006. “Plea bargains only for the guilty.” *The Journal of Law and Economics* 49(1):353–364.
- Brewer, Russell and Peter Grabosky. 2014. “The unraveling of public security in the United States: The dark side of police-community co-production.” *American Journal of Criminal Justice* 39:139–154.
- Cheng, Tony. 2023. *The Policing Machine: Enforcement, Endorsements, and the Illusion of Public Input*. University of Chicago Press.
- Clark, Tom S. 2024. “A Political-Economic Model of Police Administration.” *Journal of Politics* 86(4):1462–1478.
- Clark, Tom S. 2026. “Coordinated Justice: The Politics of Prosecution and Policing.”. Stanford University working paper.
- Clark, Tom S. Forthcoming. “Stacking the Charges: Prosecutorial Discretion in Criminal Charging Decisions.” *American Journal of Political Science* .
- Clark, Tom S., Elisha Cohen, Adam N. Glynn, Michael Leo Owens, Anna Gunderson and

- Kaylyn Jackson Schiff. 2023. "Are police racially biased in the decision to shoot?" *The Journal of Politics* 85(3):826–842.
- Clark, Tom S, Michael Leo Owens and Adam N Glynn. 2025. "Deadly force: Police shootings in urban America."
- Davis, Angela J. 2007. *Arbitrary justice: The power of the American prosecutor*. Oxford University Press.
- Easterbrook, Frank H. 1983. "Criminal procedure as a market system." *The Journal of Legal Studies* 12(2):289–332.
- Gailmard, Sean and John W. Patty. 2007. "Slackers and zealots: Civil service, policy discretion, and bureaucratic expertise." *American Journal of Political Science* 51(4):873–889.
- Garoupa, Nuno. 2012. The economics of prosecutors. In *Research Handbook on the Economics of Criminal Law*, ed. Alon Harol and Keith N. Hylton. Northampton, MA: Edward Elgar chapter 9, pp. 231–242.
- Gordon, Sanford C. 2009. "Assessing partisan bias in federal public corruption prosecutions." *American Political Science Review* 103(4):534–554.
- Gordon, Sanford C. and Gregory A. Huber. 2002. "Citizen oversight and the electoral incentives of criminal prosecutors." *American Journal of Political Science* pp. 334–351.
- Gordon, Sanford C. and Gregory A. Huber. 2009. "The political economy of prosecution." *Annual Review of Law and Social Science* 5(1):135–156.
- Hayek, Friedrich A. 1948. *Individualism and Economic Order*. Chicago: University of Chicago Press.
- Hessick, Carissa Byrne. 2021. *Punishment without trial: Why plea bargaining is a bad deal*. Abrams.

- Hübert, Ryan and Andrew T Little. 2023. “A behavioural theory of discrimination in policing.” *The Economic Journal* 133(655):2828–2843.
- Kamenica, Emir and Matthew Gentzkow. 2011. “Bayesian persuasion.” *American Economic Review* 101(6):2590–2615.
- Knox, Dean, Will Lowe and Jonathan Mummolo. 2020. “Administrative records mask racially biased policing.” *American Political Science Review* 114(3):619–637.
- Kohler-Hausmann, Issa. 2018. *Misdemeanorland: Criminal Courts and Social Control in an Age of Broken Windows Policing*. Princeton, NJ: Princeton University Press.
- Kohler, Jeremy. 2023. “Police Resistance and Politics Undercut the Authority of Prosecutors Trying to Reform the Justice System.” ProPublica.
- Landes, William M. 1971. “An economic analysis of the courts.” *The Journal of Law and Economics* 14(1):61–107.
- Lipsky, Michael. 1980. *Street-Level Bureaucracy: Dilemmas of the Individual in Public Services*. New York, NY: Russell Sage Foundation.
- Little, Andrew T and Hannah K Simpson. 2024. “Guilt and Guilty Pleas.” *American Political Science Review* pp. 1–15.
- Makris, Miltiadis. 2006. “Political authority, expertise and government bureaucracies.” *Public Choice* 127(3):267–284.
- Maynard-Moody, Steven and Michael Musheno. 2003. *Cops, Teachers, Counselors: Stories from the Front Lines of Public Service*. Ann Arbor, MI: University of Michigan Press.
- Miller, Gary J. 1977. “Bureaucratic compliance as a game on the unit square.” *Public Choice* 29(1):37–51.

- Mummolo, Jonathan. 2018. "Modern police tactics, police-citizen interactions, and the prospects for reform." *The Journal of Politics* 80(1):1–15.
- Ouss, Aurélie. 2020. "Misaligned incentives and the scale of incarceration in the United States." *Journal of Public Economics* 191:104285.
- Pennington, Nancy and Reid Hastie. 1992. "Explaining the Evidence: Tests of the Story Model for Juror Decision Making." *Journal of Personality and Social Psychology* 62(2):189–206.
- Polanyi, Michael. 1966. *The Tacit Dimension*. Chicago: University of Chicago Press.
- Rakoff, Jed S. 2016. "Why prosecutors rule the criminal justice system-And what can be done about it." *Nw. UL Rev.* 111:1429.
- Richman, Daniel. 2003. "Prosecutors and Their Agents, Agents and Their Prosecutors." *Columbia Law Review* 103(4):749–832.
- Romero, Maybell. 2019. "Prosecutors and police: An unholy union." *U. Rich. L. Rev.* 54:1097.
- Shjarback, John A., David C. Pyrooz, Scott E. Wolfe and Scott H. Decker. 2017. "Depolicing and crime in the wake of Ferguson: Racialized changes in the quantity and quality of policing among Missouri police departments." *Journal of Criminal Justice* 50:42–52.
- Starr, Sonja B. and M. Marit Rehavi. 2013. "Mandatory sentencing and racial disparity: Assessing the role of prosecutors and the effects of Booker." *Yale Law Journal* 123:2.
- Steunenberg, Bernard. 1996. "Agent discretion, regulatory policymaking, and different institutional arrangements." *Public Choice* 86(3):309–339.
- Stuart, Forrest. 2016. *Down, Out, and Under Arrest: Policing and Everyday Life in Skid Row*. Chicago, IL: University of Chicago Press.

Thompson, A.C. and Brett Murphy. 2024. “How Police and Politicians Have Undermined Reform Prosecutors.”.

URL: <https://www.propublica.org/article/police-politicians-undermined-reform-prosecutors-chicago-philadelphia>

Weingast, Barry R. 1984. “The congressional-bureaucratic system: A principal agent perspective (with applications to the SEC).” *Public Choice* 44(1):147–191.

Wilson, James Q. 1968. *Varieties of police behavior: The management of law and order in eight communities*. Harvard University Press.

Supplementary online material

Proof of Proposition 1: If the players are playing separating strategies, then upon observing $a = H$, the Prosecutor's posterior belief that the suspect is a high type is given by $\hat{p} = \Pr(\theta = H|a = H) = 1$ and upon observing $a = L$ by $\Pr(\theta = L|a = L) = 1$. If the Officer chooses $a = H$ and the Prosecutor chooses $c = L$, then she expects to receive $-\delta$. If she chooses $c = H$ then she expects to receive η . By assumption $\eta > 0$, which implies she strictly prefers $c(H) = H$. If the Officer chooses $a = L$, then she expects to receive λ from $c = L$ and $(1 - \pi)\eta - \epsilon$ from $c = H$. $\lambda > (1 - \pi)\eta - \epsilon = \underline{\lambda}$ implies she prefers $c(L) = L$. Incentive compatibility requires also that the Officer prefers pure separation. Assuming $c(L) = L$ and $c(H) = H$, if $\theta = L$ the Officer expects to receive γ from $a = L$ and $\beta - \pi\phi$ from $a = H$. $\pi \geq \bar{\pi}$ implies he prefers $a = L$. If $\theta = H$, he expects to receive $\gamma + \beta$ from $a = H$ and 0 from $a = L$. $\gamma, \beta > 0$ implies he strictly prefers $a = H$. ■

Proof of Proposition 2: In a pooling equilibrium on $a = H$, the Prosecutor's on-path posterior belief about the suspect type is equal to her prior, p . Thus, pooling requires:

$$\begin{aligned} EU_P(c = H|a = H) &> EU_P(c = L|a = H) \\ p\eta + (1 - p)(1 - \pi)\eta &> (1 - p)\lambda - \delta \\ \lambda &< \frac{p\eta + (1 - p)(1 - \pi)\eta + \delta}{1 - p} = \lambda^*. \end{aligned} \tag{1}$$

Pooling also requires that the Officer weakly prefers to choose $a = H$ in both states, for some sequentially rational response by the Prosecutor to the off-path message $a = L$. Let $r \in [0, 1]$ denote the probability the Prosecutor plays $c = H$ upon observing $a = L$. If $\theta = H$, the Officer's equilibrium payoff is $\gamma + \beta$, while deviating to $a = L$ yields $r\beta + (1 - r) \cdot 0 < \gamma + \beta$ for every r . If $\theta = L$, the Officer's equilibrium payoff is $\beta - \pi\phi$, while deviating yields $r(\beta - \pi\phi) + (1 - r)\gamma$; $\pi < \bar{\pi}$ implies $\beta - \pi\phi > \gamma$, so the deviation is unprofitable for every r . The off-path response may therefore be specified freely; in particular, the truth-consistent

belief $\Pr(\theta = L|a = L) = 1$ with response $c = L$ (sequentially rational given $\lambda > \underline{\lambda}$) supports the equilibrium, so the construction does not rely on beliefs that fail standard refinements. For the case $\pi \geq \bar{\pi}$ discussed in the footnote to the proposition: if $\lambda < \underline{\lambda}$, the Prosecutor plays $c = H$ after either message regardless of beliefs (Corollary 1), making the Officer indifferent in both states, so pooling is a best response; if $\lambda > \underline{\lambda}$, supporting pooling requires an off-path response of $c = H$, which is sequentially rational only under beliefs assigning sufficient weight to $\theta = H$ after $a = L$ —beliefs that fail the Intuitive Criterion, since only the $\theta = L$ type could profit from the deviation when $\gamma > \beta - \pi\phi$. ■

Proof of Proposition 3: If $\Pr(a = H|\theta = L) = \sigma_O$ and $\Pr(a = H|\theta = H) = 1$, then the Prosecutor’s posterior belief is given by $\Pr(\theta = H|a = H) = \frac{p}{p+(1-p)\sigma_O} \equiv \hat{p}$. To be indifferent between the two strategies in her choice space, the following condition must hold:

$$\begin{aligned}
EU_P(c = H|a = H) &= EU_P(c = L|a = H) \\
\hat{p}\eta + (1 - \hat{p})(1 - \pi)\eta &= (1 - \hat{p})\lambda - \delta \\
\hat{p} &= \frac{\lambda - (1 - \pi)\eta - \delta}{\lambda + \pi\eta} \\
\frac{p}{p + (1 - p)\sigma_O} &= \frac{\lambda - (1 - \pi)\eta - \delta}{\lambda + \pi\eta} \\
\sigma_O &= \frac{p(\delta + \eta)}{(1 - p)(\lambda - (1 - \pi)\eta - \delta)}
\end{aligned} \tag{2}$$

For σ_O to be a proper probability, two conditions must hold. The denominator must be positive, $\lambda > (1 - \pi)\eta + \delta$, and $\sigma_O \leq 1$, which rearranges to $p(\eta + \delta) \leq (1 - p)(\lambda - (1 - \pi)\eta - \delta)$, i.e., $\lambda \geq \frac{p\eta + (1 - p)(1 - \pi)\eta + \delta}{1 - p} = \lambda^*$. Since $\lambda^* - [(1 - \pi)\eta + \delta] = \frac{p(\eta + \delta)}{1 - p} > 0$, the binding condition is $\lambda \geq \lambda^*$: σ_O is a proper probability precisely when $\lambda \geq \lambda^*$, with $\sigma_O = 1$ at $\lambda = \lambda^*$ and σ_O strictly decreasing in λ thereafter. Assuming semi-separation, the Prosecutor’s posterior belief upon observing $a = L$ is $\Pr(\theta = L|a = L) = 1$. Thus, to sustain equilibrium behavior,

it must also be the case that

$$EU_P(c = L|a = L) > EU_P(c = H|a = L) \quad (3)$$

$$\lambda > (1 - \pi)\eta - \epsilon$$

$$\lambda > \underline{\lambda}$$

Because $\lambda^* - \underline{\lambda} = \frac{p\eta + \delta}{1-p} + \epsilon > 0$, this condition is implied by $\lambda \geq \lambda^*$ and imposes no further restriction. In order for the equilibrium to be incentive compatible, given a mixing probability, σ_P , for the Prosecutor upon observing $a = H$, the following condition must be satisfied as well:

$$EU_O(a = H; \sigma_P|\theta = L) = EU_O(a = L; \sigma_P|\theta = L) \quad (4)$$

$$\sigma_P(\beta - \pi\phi) - \delta(1 - \sigma_P) = \gamma$$

$$\sigma_P = \frac{\gamma + \delta}{\beta + \delta - \pi\phi}$$

It must also be the case that

$$EU_O(a = H; \sigma_P|\theta = H) \geq EU_O(a = L; \sigma_P|\theta = H) \quad (5)$$

$$\sigma_P(\gamma + \beta) - \delta(1 - \sigma_P) \geq 0$$

$$\frac{\delta}{\beta + \delta + \gamma} \leq \sigma_P = \frac{\gamma + \delta}{\beta + \delta - \pi\phi}$$

which is always true, because $\gamma, \phi > 0$ and $\pi \in (0, 1)$. It remains to confirm that σ_P is a proper probability: $\pi < \bar{\pi}$ implies $\beta - \pi\phi > \gamma$, so the denominator satisfies $\beta + \delta - \pi\phi > \gamma + \delta > 0$ and hence $\sigma_P \in (0, 1)$. ■

Proof of Proposition 4: In any uninformative equilibrium, the Officer's allegation is independent of θ , so the Prosecutor's posterior belief at any on-path message equals her prior, p . Note first the equivalence asserted in the proposition: the Prosecutor weakly

prefers $c = H$ to $c = L$ at belief p following a high-type allegation whenever

$$\begin{aligned}
EU_P(c = H|a = H) &\geq EU_P(c = L|a = H) \\
p\eta + (1 - p)(1 - \pi)\eta &\geq (1 - p)\lambda - \delta \\
p &\geq \frac{\lambda - (1 - \pi)\eta - \delta}{\lambda + \pi\eta} \equiv \bar{p},
\end{aligned} \tag{6}$$

and rearranging the middle line in λ gives $\lambda \leq \lambda^*$. Form (i) is then immediate: for $\pi < \bar{\pi}$ and $\lambda \leq \lambda^*$, the all- H profile with $c = H$ is the pooling equilibrium established in Proposition 2, and it is uninformative.

For form (ii), consider the profile in which the Officer always alleges $a = L$ and the Prosecutor plays $c = L$ after both messages. On path ($a = L$), the Prosecutor's belief is p and she prefers $c = L$ whenever $(1 - p)\lambda \geq p\eta + (1 - p)(1 - \pi)\eta - \epsilon$, i.e., $\lambda \geq \lambda^* - \frac{\delta + \epsilon}{1 - p}$. Off path ($a = H$), the response $c = L$ is sequentially rational under the belief $\Pr(\theta = H|a = H) = 0$ whenever $\lambda - \delta \geq (1 - \pi)\eta$, i.e., $\lambda \geq (1 - \pi)\eta + \delta$. Given the response $c = L$ to both messages, Officer optimality is immediate: if $\theta = H$, he receives 0 from $a = L$ and $-\delta$ from $a = H$; if $\theta = L$, he receives γ from $a = L$ and $-\delta$ from $a = H$. Both types strictly prefer $a = L$. Hence the all- L uninformative equilibrium exists if and only if $\lambda \geq \max \left\{ (1 - \pi)\eta + \delta, \lambda^* - \frac{\delta + \epsilon}{1 - p} \right\}$. Both bounds are strictly below λ^* (the first by $\frac{p(\eta + \delta)}{1 - p}$, the second by $\frac{\delta + \epsilon}{1 - p}$), so existence is guaranteed in particular whenever $\lambda \geq \lambda^*$.

Non-uniqueness follows by inspection of the parameter space. If $\lambda < \lambda^*$ and $\pi < \bar{\pi}$, the pooling equilibrium of Proposition 2 exists. If $\lambda < \lambda^*$, $\pi \geq \bar{\pi}$, and the all- L equilibrium exists, then $\lambda \geq (1 - \pi)\eta + \delta > \underline{\lambda}$, so the separating equilibrium of Proposition 1 exists. If $\lambda \geq \lambda^*$, then either $\pi < \bar{\pi}$ and the semi-separating equilibrium of Proposition 3 exists, or $\pi \geq \bar{\pi}$ and (because $\lambda \geq \lambda^* > \underline{\lambda}$) the separating equilibrium exists.

A third uninformative profile—the Officer pooling on $a = L$ met with the charge $c = H$ —can also be sustained as a perfect Bayesian equilibrium for some parameter values. It fails the Intuitive Criterion, however: only the Officer facing a high-type suspect could profit from

the off-path high-type allegation, so the Criterion requires the belief $\Pr(\theta = H|a = H) = 1$, under which the profile unravels. I therefore set it aside; the two forms in Proposition 4 are the only uninformative equilibria in which allegations carry their natural meaning.

Finally, the refinement claims in Remark 4. In the all- L equilibrium with $\pi > \bar{\pi}$, the deviation $a = H$ can strictly profit the $\theta = H$ type (who would gain $\gamma + \beta > 0$ if met with $c = H$) but cannot profit the $\theta = L$ type under any response (his best deviation payoff, $\beta - \pi\phi$, is strictly less than his equilibrium payoff γ when $\pi > \bar{\pi}$). The Intuitive Criterion therefore requires the belief $\Pr(\theta = H|a = H) = 1$, to which the unique best response is $c = H$, and the $\theta = H$ type then strictly profits by deviating; the equilibrium fails the Criterion. For $\pi < \bar{\pi}$, both types strictly profit from $a = H$ when met with $c = H$ ($\gamma + \beta > 0$ and $\beta - \pi\phi > \gamma$, respectively), so the Criterion places no restriction on the off-path belief and the all- L equilibrium survives. ■

Robustness: a symmetric trust cost

In the baseline model the trust cost δ is paid only when the Prosecutor downgrades a high-type allegation ($a = H, c = L$). Here I show that symmetrizing this assumption—charging δ to both players on *any* misalignment between allegation and charge, so that escalation ($a = L, c = H$) is also penalized—leaves every headline result intact and, if anything, enlarges the informative and standing-down regions. Formally, subtract δ from both players in the two ($a = L, c = H$) cells of Table 1.

Corollary 2 *Suppose the trust cost δ is paid on any misalignment between allegation and charge. Then: (i) the pooling boundary λ^* (Definition 3) and the semi-separating mixing probabilities σ_O^* and σ_P^* (Proposition 3) are unchanged; (ii) the escalation threshold falls from $\underline{\lambda}$ to $\underline{\lambda} - \delta$, so the separating equilibrium (Proposition 1) is sustained on a weakly larger range of λ ; and (iii) the all-low “standing down” equilibrium (Proposition 4(ii)) exists whenever $\lambda \geq \max \left\{ (1 - \pi)\eta + \delta, \lambda^* - \frac{2\delta + \epsilon}{1 - p} \right\}$, its on-path bound falling by $\frac{\delta}{1 - p}$, so it too is sustained on a weakly larger range of λ . In particular, the observational-equivalence result (Result 1) is unaffected: symmetrizing the trust cost introduces no new equilibrium forms and does not alter their charge profiles; it weakly expands the parameter regions in which the informative and standing-down equilibria exist.*

Proof The Prosecutor’s on-path comparison after a high-type allegation— $c = H$ versus the downgrade $c = L$ —involves only the $(a = H, \cdot)$ cells, which the symmetric change does not touch; hence the boundary λ^* and the posterior $\hat{p}$ that pins σ_O^* are identical to the baseline. The Officer’s indifference that pins σ_P^* compares $a = H$ (met by the Prosecutor’s mixture over c) with $a = L$ met by $c = L$ on path; neither side uses the $(a = L, c = H)$ cell, so σ_P^* is unchanged. This gives (i). For (ii), in the separating equilibrium the Prosecutor, having observed $a = L$ and inferred $\theta = L$, compares $c = L$ (payoff λ) with escalation to $c = H$ (payoff $(1 - \pi)\eta - \epsilon - \delta$ under the symmetric cost); she prefers $c = L$ iff $\lambda > (1 - \pi)\eta - \epsilon - \delta = \underline{\lambda} - \delta$, weaker than the baseline condition $\lambda > \underline{\lambda}$, and all other constraints in the proof of Proposition 1 are unchanged. For (iii), in the all-low equilibrium the Officer pools on $a = L$, so the Prosecutor’s on-path belief is her prior p , and her willingness to play $c = L$ rather than escalate requires $(1 - p)\lambda \geq p\eta + (1 - p)(1 - \pi)\eta - \epsilon - \delta$, i.e. $\lambda \geq \lambda^* - \frac{2\delta + \epsilon}{1 - p}$, which is $\frac{\delta}{1 - p}$ below the baseline bound; the off-path constraint requiring $\lambda \geq (1 - \pi)\eta + \delta$ uses only the downgrade cell and is unchanged, and Officer optimality is immediate as in Proposition 4. Because the equilibrium set and the charge profiles ($c^* = a^*$ in separation, pooling, and both uninformative forms) are unchanged, Result 1 carries over verbatim. ■

Intuitively, penalizing escalation as well as downgrading makes the Prosecutor more reluctant to move *away* from the Officer’s allegation in either direction. It has no effect where the binding tension is her temptation to downgrade—which governs λ^* , σ_O^* , and σ_P^* —and it makes her weakly more willing to accept low-type allegations, expanding both the separating region (she less often escalates a known low type) and the standing-down region (she less often escalates a pooled low allegation). The directional trust cost of the baseline is therefore a conservative choice: it does not generate the observational-equivalence result or the degradation of information, and relaxing it only strengthens the informative and standing-down equilibria.

Robustness: the Officer's taste for severity

The main text notes that the model's central tension requires the Officer to lean punitive. The following corollaries formalize the sense in which the results depend on a taste for severity that is both positive and large relative to the justice motive γ .

Corollary 3 *If the Officer's taste for severity is bounded by his justice motive, $-\gamma < \beta < \gamma$, then $\bar{\pi} < 0$, and for every court disposition $\pi \in (0, 1)$: (i) no pooling (Proposition 2) or semi-separating (Proposition 3) equilibrium exists, each requiring $\pi < \bar{\pi}$; and (ii) provided $\lambda > \underline{\lambda}$, the separating equilibrium of Proposition 1 exists. The Officer's allegations can be truthful regardless of the court; where $\lambda < \underline{\lambda}$, however, the Prosecutor over-charges known low types (Corollary 1), so informative allegations need not yield type-matched charges.*

Proof $-\gamma < \beta < \gamma$ gives $\bar{\pi} = \frac{\beta - \gamma}{\phi} < 0 < \pi$ for all admissible π , so the hypotheses of Propositions 2 and 3 are never met; consistently, $\sigma_P^* = \frac{\gamma + \delta}{\beta + \delta - \pi\phi}$ exceeds one whenever $\pi > \bar{\pi}$, so no interior semi-separating mixture exists. For separation, the Officer facing a low-type suspect prefers a truthful $a = L$ (payoff γ) to $a = H$ met by $c = H$ (payoff $\beta - \pi\phi < \gamma$ when $\pi > \bar{\pi}$), and the Officer facing a high-type suspect prefers a truthful $a = H$ (payoff $\gamma + \beta > 0$, using $\beta > -\gamma$) to $a = L$ (payoff 0); the Prosecutor adopts the allegations when $\lambda > \underline{\lambda}$ (Proposition 1). The over-charging claim is Corollary 1. ■

Corollary 4 *If the Officer has no taste for severity, $\beta \leq 0$, his motive to misrepresent dissolves: the Officer facing a low-type suspect never strictly prefers that the suspect be charged high. Any over-charging is then driven solely by the Prosecutor, and arises only where $\lambda < \underline{\lambda}$ (Corollary 1). If in addition $\beta > -\gamma$, separation obtains for $\lambda > \underline{\lambda}$ and any court; if $\beta \leq -\gamma$, the Officer facing a high-type suspect prefers to under-allege, and natural-meaning separation is not incentive compatible.*

Proof For $\theta = L$, the Officer earns $\gamma > 0$ from a truthful $a = L$ adopted as $c = L$, while any high charge yields $\beta - \pi\phi \leq -\pi\phi < 0$; he therefore never strictly prefers to induce a high charge, and over-charging arises only when the Prosecutor herself prefers $c = H$ against a known low type, i.e. when $\lambda < \underline{\lambda}$ (Corollary 1). For $\theta = H$, a truthful $a = H$ adopted as $c = H$ yields $\gamma + \beta$, which exceeds the deviation payoff 0 only when $\beta > -\gamma$; for $\beta \leq -\gamma$ the

high type prefers $a = L$, so the separating profile fails incentive compatibility. ■

These corollaries locate the interesting region of the parameter space: it is precisely when the Officer's taste for severity is positive and large enough ($\beta > \gamma > 0$, so $\bar{\pi} > 0$) that the court's disposition becomes pivotal and the pooling, semi-separating, and standing-down equilibria emerge. This is the case on which the analysis in the main text focuses.